

SKORPIO: Advanced Binary Instrumentation Framework

NGUYEN Anh Quynh <aquynh -at- gmail.com>



About me

- **Nguyen Anh Quynh**, aquynh-at-gmail.com
 - ▶ Nanyang Technological University, Singapore
 - ▶ PhD in Computer Science
 - ▶ Operating System, Virtual Machine, Binary analysis, etc
 - ▶ Usenix, ACM, IEEE, LNCS, etc
 - ▶ Blackhat USA/EU/Asia, DEFCON, Recon, HackInTheBox, Syscan, etc
 - ▶ Capstone disassembler: <http://capstone-engine.org>
 - ▶ Unicorn emulator: <http://unicorn-engine.org>
 - ▶ Keystone assembler: <http://keystone-engine.org>



Agenda

- 1 Dynamic Binary Instrumentation (DBI)
- 2 Skorpio instrumentation engine
- 3 Demos
- 4 Conclusions

Dynamic Binary Instrumentation (DBI)


Definition

- A method of analyzing a binary application at runtime through injection of instrumentation code.
 - ▶ Extra code executed as a part of original instruction stream
 - ▶ No change to the original behavior
- Framework to build apps on top of it

Applications

- Code tracing/logging
- Debugging
- Profiling
- Security enhancement/mitigation

DBI illustration




DBI techniques

- Just-in-Time translation
 - ▶ Transparently translate & execute code at runtime
 - ★ Perform on IR: Valgrind
 - ★ Perform directly on native code: DynamoRio
 - ▶ Better control on code executed
 - ▶ Heavy, super complicated in design & implementation
- Hooking
 - ▶ Lightweight, much simpler to design & implement
 - ▶ Less control on code executed & need to know in advance where to instrument

Hooking mechanisms - Inline

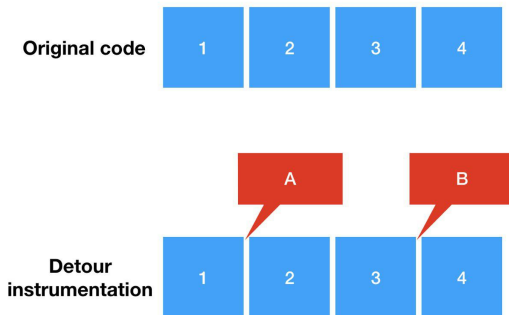
- Inline code injection

- ▶ Put instrumented code inline with original code
- ▶ Can instrument anywhere & unlimited in extra code injected
- ▶ Require complicated code rewrite



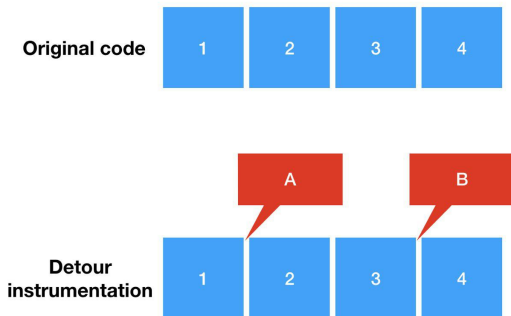
Hooking mechanisms - Detour

- Detour injection
 - ▶ Branch to external instrumentation code
 - ★ User-defined **CALLBACK** as instrumented code
 - ★ **TRAMPOLINE** memory as a step-stone buffer
 - ▶ Limited on where to hook
 - ★ Basic block too small?
 - ▶ Easier to design & implement

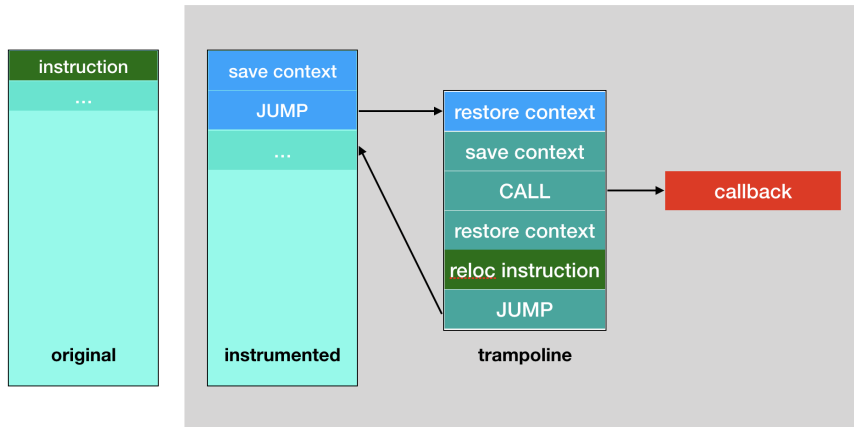


Detour injection mechanisms

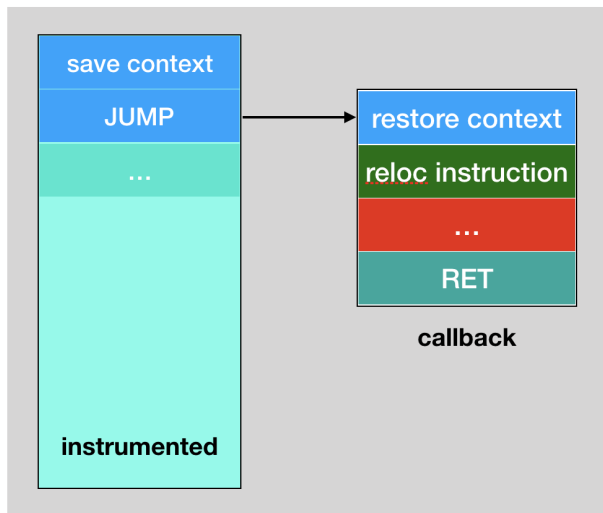

- Branch from original instruction to instrumented code
- Branch to trampoline, or directly to callback
 - ▶ Jump-trampoline technique
 - ▶ Jump-callback technique
 - ▶ Call-trampoline technique
 - ▶ Call-callback technique



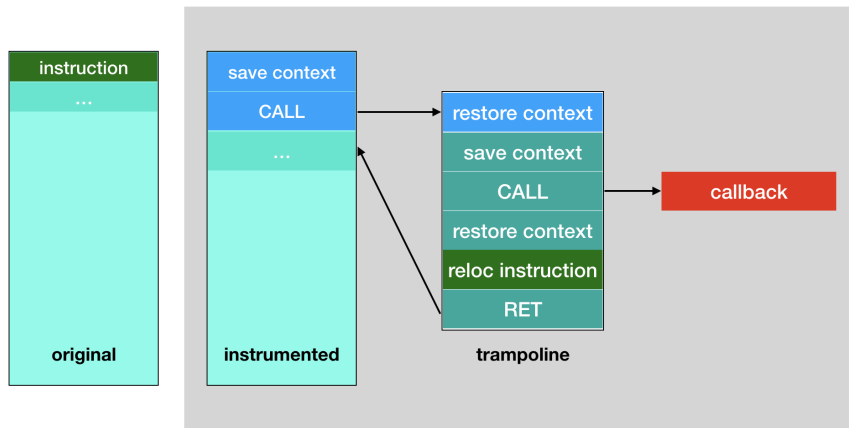
Jump-trampoline technique



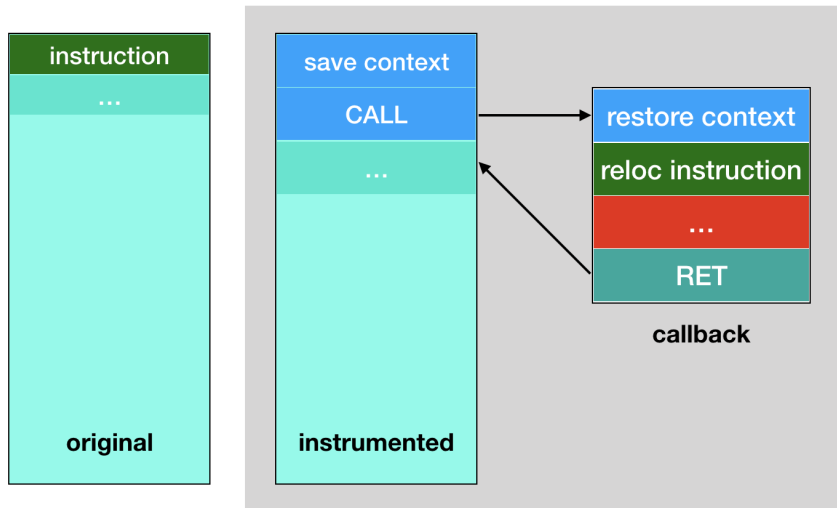
Jump-callback technique



Call-trampoline technique



Call-callback technique




Problems of existing DBI

- Limited on platform support
- Limited on architecture support
- Limited on instrumentation techniques
- Limited on optimization

SKORPIO framework

- Open source, with permissive license
- Low level framework to build applications on top
 - ▶ App typically designed as dynamic libraries (DLL/SO/DYLIB)
- Cross-platform-architecture
 - ▶ Windows, MacOS, Linux, BSD, etc
 - ▶ X86, Arm, Arm64, Mips, Sparc, PowerPC
- Allow all kind of instrumentations
 - ▶ Arbitrary address, in any privilege level
- Designed to be easy to use, but support all kind of optimization
 - ▶ Super fast (100x) compared to other frameworks, with proper setup
- Support static instrumentation, too!

SKORPIO architecture



Cross platform - Memory

- Thin layer to abstract away platform details
- Different OS supported in separate plugin
 - ▶ Posix vs Windows
- Trampoline buffer
 - ▶ Allocate memory: `malloc()` vs `VirtualAlloc()`
 - ▶ Memory privilege RWX: `mprotect()` vs `VirtualAlloc()`
 - ▶ Trampoline buffer as close as possible to code to reduce branch distance
- Patch code in memory
 - ▶ Unprotect -> Patch -> Re-protect
 - ▶ `mprotect()` vs `VirtualProtect()`

Cross architecture - Save/Restore context

- Save memory/registers modified by initial branch & callback
- Keep the code size as small as possible
- Depend on architecture + mode
 - ▶ X86-32: PUSHAD; PUSHFD & POPFD; POPAD
 - ▶ X86-64 & other CPUs: no simple instruction to save all registers :-(
 - ★ Calling convention: cdecl, optlink, pascal, stdcall, fastcall, safecall, thiscall, vectorcall, Borland, Watcom
 - ★ SystemV ABI vs Windows ABI
- Special API to customize code to save/restore context

Cross architecture - Callback argument

- Pass user argument to user-defined callback
- Depend on architecture + mode & calling convention
 - ▶ SysV/Windows x86-32 vs x86-64
 - ★ Windows: cdecl, optlink, pascal, stdcall, fastcall, safecall, thiscall, vectorcall, Borland, Watcom
 - ▶ X86-64: "mov rcx, <value>" or "mov rdi, <value>". Encoding depends on data value
 - ▶ Arm: "ldr r0, [pc, 0]; b .+8; <4-byte-value>"
 - ▶ Arm64: "movz x0, <lo16>; movk x0, <hi16>, lsl 16"
 - ▶ Mips: "li \$a0, <value>"
 - ▶ PPC: "lis %r3, <hi16>; ori %r3, %r3, <lo16>"

Cross architecture - Branch distance

- Distance from hooking place to callback cause nightmare :-(
 - Some architectures have no explicit support for far branching
 - X86-64 JUMP: "push <addr>; ret" or "push 0; mov dword ptr [rsp+4], <addr>" or "jmp [rip]"
 - X86-64 CALL: "push <next-addr>; push <target>; ret"
 - Arm JUMP: "b <addr>" or "ldr pc, [pc, #-4]"
 - Arm CALL: "bl <addr>" or "add lr, pc, #4; ldr pc, [pc, #-4]"
 - Arm64 JUMP: "b <addr>" or "ldr x16, .+8; br x16"
 - Arm64 CALL: "bl <addr>" or "ldr x16, .+12; blr x16; b .+12"
 - Mips JUMP: "li \$t0, <addr>; jr \$t0"
 - Mips CALL: "li \$t0, <addr>; move \$t9, \$t0; jalr \$t0"
 - Sparc JUMP: "set <addr>, %l4; jmp %l4; nop"
 - Sparc CALL: "set <addr>, %l4; call %l4; nop"

Cross architecture - Branch for PPC

- PPC has no far jump instruction :-(
 - ▶ copy LR to r23, save target address to r24, then copy to LR for BLR
 - ▶ restore LR from r23 after jumping back from trampoline
 - ▶ "mflr %r23; lis %r24, <hi16>; ori %r24, %r24, <lo16>; mtlr %r24; blr"
- PPC has no far call instruction :-(
 - ▶ save r24 with target address, then copy r24 to LR
 - ▶ point r24 to instruction after BLR, so later BLR go back there from callback
 - ▶ "lis %r24, <target-hi16>; ori %r24, %r24, <target-lo16>; mtlr %r24; lis %r24, <ret-hi16>; ori %r24, %r24, <ret-lo16>; blr"

```
SK_INLINE_NO static void bbb_hook(size_t v)
{
    // restore LR from R24
    __asm__("mtlr %r24");

    printf("== in callback, userdata = %zu\n", v);

    return;
}
```

Cross architecture - Scratch register

- Scratch registers used in initial branching
 - ▶ Arm64, Mips, Sparc & PPC do not allow branch to indirect target in memory
 - ▶ Calculate branch target, or used as branch target
 - ▶ Need scratch register(s) that are unused in local context
 - ★ Specified by user via API, or discovered automatically by engine

Cross architecture - Flush code cache

- Code patching need to be reflected in i-cache
- Depend on architecture
 - ▶ X86: no need
 - ▶ Arm, Arm64, Mips, PowrPC, Sparc: special syscalls/instructions to flush/invalidate i-cache
 - ▶ Linux/GCC has special function: `cacheflush(begin, end)`

Code boudary & relocation

- Need to extract instructions overwritten at instrumentation point
 - ▶ Determine instruction boundary for X86
 - ▶ Use Capstone disassembler
- Need to rewrite instructions to work at relocated place (trampoline)
 - ▶ Relative instructions (branch, memory access)
 - ▶ Use Capstone disassembler to detect instruction type
 - ▶ Use Keystone assembler to recompile



Code analysis

- Avoid overflow to next basic block
 - ▶ Analysis to detect if basic block is too small for patching
- Reduce number of registers saved before callback
- Registers to be chosen as scratch registers

Customize on instrumentation

- API to setup calling convention
- User-defined callback
- User-defined trampoline
- User-defined scratch registers
- User-defined save-restore context
- User-defined code to setup callback args
- Patch hooks in batch, or individual
- User decide when to write/unwrite memory protect

Skorpio sample C code

Sample for Skorpio engine

--- Original code

BBB code = 0x400ca0, callback = 0x400c80

Hook info:

Hook type: 2

Hook address: 0x400ca0

Hook callback: 0x400c80

Hook user_data: 0x7b

Hook trampoline addr: 0x7f1aa7911000

Hook trampoline size: 86

Hook trampoline code: 5053515257565541504151415241549c48c7c77b0000006a00c70424321091a7c74424041a7f00006a00c70424800c4000c39d415c415a415941585d5e5f5a595b584883ec08b9800c4000baa00c400068ae0c4000c3

Patch size: 14

Patched code: ff250000000001091a71a7f0000

Hook original code size: 14

Hook original code: 4883ec08b9800c4000baa00c4000

--- Functions with instrumentation now

== inside callback, userdata = 123

BBB code = 0x400ca0, callback = 0x400c80

--- Restored original code, now without instrumentation

BBB code = 0x400ca0, callback = 0x400c80

Status

- Cross-platforms: Windows, Linux, MacOS, BSD
- Python binding available
- Need to test on Android & iOS
- Cross-architecture: X86, Arm, Arm64, Mips, PowerPC, Sparc
- More test before public release - soon

Conclusions

- **SKORPIO** is an advanced framework for binary instrumentation
 - ▶ Open-source, cross-platform-architecture
 - ▶ All level of customization for better performance
 - ▶ Dynamic & static instrumentation
 - ▶ Lay the foundation for future security tools R&D



Acknowledgement

- Demo on Darko fuzzer is co-worked with Dr.Wei Lei (NTU)
- Huge thanks to @capstone_engine & @keystone_engine communities for great support!
- @_hugsy_ for Qemu images of Mips, PowerPC & Sparc

SKORPIO: Advanced Binary Instrumentation Framework

NGUYEN Anh Quynh <aquynh -at- gmail.com>

